

# Final Project Submission

Please fill out:

- Student name: Justin Grisanti
- Student pace: self-paced
- Scheduled project review date/time: 1/5/2022 @ TBD
- Instructor name: Claude Fried
- Blog post URL: [https://justingrisanti.github.io/predicting\\_rain\\_patterns\\_in\\_australia](https://justingrisanti.github.io/predicting_rain_patterns_in_australia)

Total Time to Run: approx. 27 minutes

## Section 1: Business Understanding

The purpose of this section is to define the business problem and understand the stakeholders for the work that I am performing. The Bureau of Meteorology is responsible for predicting weather patterns throughout the entire Australian region. According to their website, their forecast accuracy for rain varies much more than their forecasts for temperature and wind.

According to their analyses, they've underpredicted rainfall each year for the past five years. The goal is to create a classification model that allows the Bureau of Meteorology to improve their predictions of whether or not it will rain the next day. This will allow them to inform the public better so that citizens can prepare accordingly for the possibility of rain.

The stakeholders of this project are the Bureau of Meteorology and citizens of Australia.

The main purpose of this classification model is predictive, meaning that given characteristics of rain data on a given day, the model should be able to predict whether it will rain the next day or not. My model is not meant to replace the Bureau of Meteorology's current system of predicting rain for the region of Australia, however, it is meant serve as an input to strengthen their predictions and assumptions, and to reduce the risk of failing to predict that it will rain the next day.

## Section 2: Data Understanding

After scanning the data, we have weather-related information for a period of 12/1/2008 to 6/25/2017—8 years, 6 months and 24 days worth of data.

As shown below, this data has many different weather-related metrics, such as the wind speed, humidity, pressure, whether it was sunny or cloudy, and temperature. These seem like appropriate parameters to run a classification-based model in order to predict whether or not it will rain the next day, and I do not sense any limitations from this data.

```
In [1]: # Import Libraries

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
from sklearn import tree
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.impute import SimpleImputer
from sklearn.neighbors import KNeighborsClassifier
from sklearn.pipeline import Pipeline
from sklearn.model_selection import train_test_split, GridSearchCV, cross_v
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier
from sklearn.metrics import accuracy_score
from sklearn.impute import SimpleImputer
from sklearn.metrics import make_scorer, classification_report, log_loss, f
from sklearn.model_selection import StratifiedKFold, cross_val_score
from sklearn.base import clone
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
import statsmodels as sm
import sklearn.preprocessing as preprocessing
from scipy import stats
import seaborn as sns
from imblearn.over_sampling import SMOTE

import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: # Import data from csv

rain_data = pd.read_csv('data/WeatherAUS.csv')
```

```
In [3]: rain_data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 145460 entries, 0 to 145459
Data columns (total 23 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                  145460 non-null object
1   Location              145460 non-null object
2   MinTemp               143975 non-null float64
3   MaxTemp               144199 non-null float64
4   Rainfall              142199 non-null float64
5   Evaporation           82670 non-null  float64
6   Sunshine              75625 non-null  float64
7   WindGustDir           135134 non-null object
8   WindGustSpeed         135197 non-null float64
9   WindDir9am            134894 non-null object
10  WindDir3pm            141232 non-null object
11  WindSpeed9am          143693 non-null float64
12  WindSpeed3pm          142398 non-null float64
13  Humidity9am           142806 non-null float64
14  Humidity3pm           140953 non-null float64
15  Pressure9am           130395 non-null float64
16  Pressure3pm           130432 non-null float64
17  Cloud9am              89572 non-null  float64
18  Cloud3pm              86102 non-null  float64
19  Temp9am               143693 non-null float64
20  Temp3pm               141851 non-null float64
21  RainToday             142199 non-null object
22  RainTomorrow          142193 non-null object
dtypes: float64(16), object(7)
memory usage: 25.5+ MB

```

Please see the following column descriptions:

**MinTemp:** The minimum temperature in degrees celsius

**MaxTemp:** The maximum temperature in degrees celsius

**Rainfall:** The amount of rainfall recorded for the day in mm

**Evaporation:** The so-called Class A pan evaporation (mm) in the 24 hours to 9am

**Sunshine:** The number of hours of bright sunshine in the day.

**WindGustDir:** The direction of the strongest wind gust in the 24 hours to midnight

**WindGustSpeed:** The speed (km/h) of the strongest wind gust in the 24 hours to midnight

**WindDir9am:** Direction of the wind at 9am

**WindDir3pm:** Direction of the wind at 3pm

**WindSpeed9am:** Wind speed (km/hr) averaged over 10 minutes prior to 9am

**WindSpeed3pm:** Wind speed (km/hr) averaged over 10 minutes prior to 3pm

**Humidity9am:** Humidity (percent) at 9am

**Humidity3pm:** Humidity (percent) at 3pm

**Pressure9am:** Atmospheric pressure (hpa) reduced to mean sea level at 9am

**Pressure3pm:** Atmospheric pressure (hpa) reduced to mean sea level at 3pm

**Cloud9am:** Fraction of sky obscured by cloud at 9am. This is measured in "oktas", which are a unit of eighths. It records how many eighths of the sky are obscured by cloud. A 0 measure indicates completely clear sky whilst an 8 indicates that it is completely overcast.

**Cloud3pm:** Fraction of sky obscured by cloud (in "oktas": eighths) at 3pm. See Cloud9am for a description of the values

**Temp9am:** Temperature (degrees C) at 9am

**Temp3pm:** Temperature (degrees C) at 3pm

**RainToday:** Boolean: 1 if precipitation (mm) in the 24 hours to 9am exceeds 1mm, otherwise 0

**RainTomorrow:** The amount of next day rain in mm. Used to create response variable RainTomorrow. A kind of measure of the "risk".

In [4]:

```
# To show all columns and rows

pd.set_option('display.max_columns', None)
pd.set_option('display.max_rows', None)

rain_data.head()
```

Out [4]:

	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustDir	Wind
<b>0</b>	2008-12-01	Albury	13.4	22.9	0.6	NaN	NaN	W	
<b>1</b>	2008-12-02	Albury	7.4	25.1	0.0	NaN	NaN	WNW	
<b>2</b>	2008-12-03	Albury	12.9	25.7	0.0	NaN	NaN	WSW	
<b>3</b>	2008-12-04	Albury	9.2	28.0	0.0	NaN	NaN	NE	
<b>4</b>	2008-12-05	Albury	17.5	32.3	1.0	NaN	NaN	W	

In [5]:

```
rain_data[rain_data['RainTomorrow'].isna()].head()
```

Out [5]:

	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustDir	Wind
<b>14</b>	2008-12-15	Albury	8.4	24.6	0.0	NaN	NaN	NaN	
<b>283</b>	2009-09-10	Albury	2.6	NaN	0.0	NaN	NaN	NaN	
<b>435</b>	2010-02-09	Albury	22.1	35.1	0.0	NaN	NaN	NaN	
<b>437</b>	2010-02-11	Albury	21.5	35.0	0.0	NaN	NaN	NaN	
<b>443</b>	2010-02-17	Albury	15.5	30.6	0.0	NaN	NaN	NaN	

In [6]:

```
rain_data = rain_data.dropna(axis=0, subset = ['RainTomorrow'])
```

In [7]:

```
rain_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 142193 entries, 0 to 145458
Data columns (total 23 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                  142193 non-null object
1   Location              142193 non-null object
2   MinTemp               141556 non-null float64
3   MaxTemp               141871 non-null float64
4   Rainfall              140787 non-null float64
5   Evaporation           81350 non-null float64
6   Sunshine              74377 non-null float64
7   WindGustDir           132863 non-null object
8   WindGustSpeed         132923 non-null float64
9   WindDir9am            132180 non-null object
10  WindDir3pm            138415 non-null object
11  WindSpeed9am          140845 non-null float64
12  WindSpeed3pm          139563 non-null float64
13  Humidity9am           140419 non-null float64
14  Humidity3pm           138583 non-null float64
15  Pressure9am           128179 non-null float64
16  Pressure3pm           128212 non-null float64
17  Cloud9am              88536 non-null float64
18  Cloud3pm              85099 non-null float64
19  Temp9am               141289 non-null float64
20  Temp3pm               139467 non-null float64
21  RainToday             140787 non-null object
22  RainTomorrow          142193 non-null object
dtypes: float64(16), object(7)
memory usage: 26.0+ MB
```

```
In [8]: # Separate target variable from data and complete train-test split

y = rain_data['RainTomorrow']
X = rain_data.drop('RainTomorrow', axis=1)

X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=42)
```

```
In [9]: # Reset index and drop it

X_train.reset_index(inplace=True)
X_train = X_train.drop(columns=['index'], axis=1)
X_train.head()
```

Out [9]:

	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustDir	V
0	2009-04-12	Woomera	14.9	30.3	0.0	7.4	10.9	S	
1	2014-12-08	Witchcliffe	14.6	21.5	0.2	NaN	NaN	SSE	
2	2015-06-06	SalmonGums	9.0	23.7	0.0	NaN	NaN	W	
3	2014-01-09	Albany	15.3	24.0	0.0	8.2	12.1	NaN	
4	2014-12-14	Mildura	17.3	37.5	0.0	8.6	11.4	N	

In [10]:

```
# Reset index and drop it

X_test.reset_index(inplace=True)
X_test = X_test.drop(columns=['index'],axis=1)
X_test.head()
```

Out [10]:

	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustDir	V
0	2016-06-09	Ballarat	7.1	13.0	8.8	NaN	NaN	N	
1	2009-10-24	Walpole	13.2	18.3	0.0	NaN	NaN	E	
2	2015-09-21	PerthAirport	9.2	22.7	0.0	5.0	11.1	ENE	
3	2011-12-06	Cobar	15.3	26.1	0.0	10.4	NaN	E	
4	2014-03-15	Sale	11.9	31.8	0.0	5.0	4.1	NW	

In [11]:

```
y_train = y_train.replace('Yes', 1.0)
y_train = y_train.replace('No', 0.0)
y_test = y_test.replace('Yes', 1.0)
y_test = y_test.replace('No', 0.0)
```

```
In [12]: # Inspect Target Variable  
  
y_train.value_counts(normalize=True)
```

```
Out[12]: 0.0    0.775412  
        1.0    0.224588  
        Name: RainTomorrow, dtype: float64
```


```
In [13]: # Ensure we get a fair spread of data across the country  
  
X_train['Location'].value_counts()
```



```
Out[13]: Canberra      2578
          Sydney        2429
          Hobart         2416
          Perth          2408
          Brisbane       2401
          Darwin         2382
          Adelaide       2329
          Tuggeranong    2298
          Mildura        2295
          Launceston     2279
          Bendigo        2279
          PerthAirport   2269
          Woomera        2266
          Ballarat       2262
          Albany         2255
          MountGambier   2254
          Townsville     2248
          CoffsHarbour   2247
          MelbourneAirport 2246
          Watsonia       2243
          Sale           2241
          GoldCoast      2240
          Witchcliffe    2240
          AliceSprings   2237
          Portland       2236
          NorfolkIsland  2236
          WaggaWagga     2236
          Cobar          2236
          Albury         2234
          Penrith        2233
          BadgerysCreek  2232
          Cairns         2232
          Newcastle     2229
          Nuriootpa      2224
          SydneyAirport  2224
          Wollongong     2213
          SalmonGums     2206
          Richmond       2201
          Dartmoor       2200
          MountGinini    2192
          NorahHead      2179
          Moree          2134
          Walpole        2133
          PearceRAAF     2087
          Williamtown    1905
          Melbourne      1834
          Katherine      1193
          Uluru          1138
          Nhil           1135
          Name: Location, dtype: int64
```

In [14]:

```
# Plot each column in a histogram to see what type of distribution there is
columns_with_nulls = X_train.drop(['Location', 'Date'], axis=1)
columns_with_nulls.hist(figsize=(20,20))
plt.savefig('Visualizations/ColumnsHist.png', bbox_inches = 'tight')
```



In [15]:

```
X_train.info(1, null_counts=True)
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 106644 entries, 0 to 106643
Data columns (total 22 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Date                  106644 non-null object
 1   Location              106644 non-null object
 2   MinTemp               106150 non-null float64
 3   MaxTemp               106387 non-null float64
 4   Rainfall              105540 non-null float64
 5   Evaporation           60909 non-null float64
 6   Sunshine              55677 non-null float64
 7   WindGustDir           99665 non-null object
 8   WindGustSpeed         99715 non-null float64
 9   WindDir9am            99120 non-null object
10   WindDir3pm            103822 non-null object
11   WindSpeed9am          105646 non-null float64
12   WindSpeed3pm          104681 non-null float64
13   Humidity9am           105326 non-null float64
14   Humidity3pm           103917 non-null float64
15   Pressure9am           96098 non-null float64
16   Pressure3pm           96123 non-null float64
17   Cloud9am              66285 non-null float64
18   Cloud3pm              63757 non-null float64
19   Temp9am               105967 non-null float64
20   Temp3pm               104579 non-null float64
21   RainToday             105540 non-null object
dtypes: float64(16), object(6)
memory usage: 17.9+ MB

```

As we can see, Evaporation, Sunshine, and cloud data all have a large amount of nulls. Let's take a deeper look into these variables.

In [16]: `x_train.head()`

```

Out[16]:
   Date      Location  MinTemp  MaxTemp  Rainfall  Evaporation  Sunshine  WindGustDir  \
0  2009-04-12    Woomera     14.9     30.3        0.0           7.4         10.9           S
1  2014-12-08  Witchcliffe     14.6     21.5        0.2          NaN          NaN          SSE
2  2015-06-06  SalmonGums      9.0     23.7        0.0          NaN          NaN           W
3  2014-01-09     Albany     15.3     24.0        0.0           8.2         12.1          NaN
4  2014-12-14    Mildura     17.3     37.5        0.0           8.6         11.4           N

```

The following code is to create meaningful information about our null values, so that we can impute them in a more educated manner.

```
In [17]: # Dropping NAs to obtain distribution for data that isn't NA

rain_data_cloud_dropna = X_train.dropna(axis=0, subset=['Cloud3pm', 'Cloud9a
```

```
In [18]: # Dropping NAs to obtain distribution for data that isn't NA

rain_data_sunshine_dropna = X_train.dropna(axis=0, subset=['Sunshine'])
```

```
In [19]: # Checking null sunshine data against rain data to see what the population

rain_data_sunshine_nulls = X_train[X_train['Sunshine'].isna()]
rain_data_sunshine_nulls['RainToday'].value_counts()
```

```
Out[19]: No      39031
        Yes     11143
        Name: RainToday, dtype: int64
```

```
In [20]: # Removing nulls from sunshine data and biforcating the population to get m

sunshine_when_rain = rain_data_sunshine_dropna[rain_data_sunshine_dropna['R
sunshine_no_rain = rain_data_sunshine_dropna[rain_data_sunshine_dropna['Rai
```

```
In [21]: # Checking evaporation mean when sunshine is greater than zero

evaporation_test = X_train.loc[X_train['Sunshine']>0.0, 'Evaporation']
```

```
In [22]: # Checking the mean of cloud data when it doesn't rain and humidity is high
# (Population is not normal without humidity check)

cloud9_no_rain_lower_humidity = X_train.loc[(rain_data['RainToday']=='No')
cloud9_no_rain_higher_humidity = X_train.loc[(X_train['RainToday']=='No') &
cloud3_no_rain_lower_humidity = X_train.loc[(X_train['RainToday']=='No') &
cloud3_no_rain_higher_humidity = X_train.loc[(X_train['RainToday']=='No') &
```

```
In [23]: # Checking the mean of cloud data when it does rain

cloud9_rain = X_train.loc[X_train['RainToday']=='Yes', 'Cloud9am']
cloud3_rain = X_train.loc[X_train['RainToday']=='Yes', 'Cloud3pm']
```

In [24]:

```
# Checking null data to see how many records have both null cloud data and

test = []

for index in range(0,106643,1):
    if pd.isna(X_train['Cloud3pm'].loc[index]) and pd.isna(X_train['Cloud9a
        if pd.notna(X_train['Sunshine'].loc[index]):
            test.append('Sunshine')
        elif pd.isna(X_train['Sunshine'].loc[index]):
            test.append('Neither')
        else:
            pass
    elif pd.notna(X_train['Cloud3pm'].loc[index]) or pd.notna(X_train['Clou
        if pd.notna(X_train['Sunshine'].loc[index]):
            test.append('Both')
        elif pd.isna(X_train['Sunshine'].loc[index]):
            test.append('Clouds')
        else:
            pass
    else:
        pass
```

In [25]:

```
print("There are " + str(test.count('Sunshine')) + " records of sunshine da
print("There are " + str(test.count('Clouds')) + " records of cloud data wi
print("There are " + str(test.count('Neither')) + " records of neither suns
print("There are " + str(test.count('Both')) + " records of both sunshine a
```

There are 5580 records of sunshine data with no cloud data.  
 There are 19220 records of cloud data with no sunshine data.  
 There are 31746 records of neither sunshine or cloud data.  
 There are 50097 records of both sunshine and cloud data.

In [26]:

```
test_data = {'Sunshine data with Cloud Nulls': 8258, 'Cloud Data with Sunsh
```

```
In [27]: # Checking the records to see the relationship between cloud data and sunsh

test2 = []

for index in range(0,106643,1):
    if X_train['Cloud3pm'].loc[index]== 0 or X_train['Cloud9am'].loc[index]
        if X_train['Sunshine'].loc[index] != 0:
            test2.append('Sunshine')
        elif X_train['Sunshine'].loc[index]==0:
            test2.append('Neither')
        else:
            pass
    elif X_train['Cloud3pm'].loc[index]!=0 or X_train['Cloud9am'].loc[index]
        if X_train['Sunshine'].loc[index]!=0:
            test2.append('Both')
        elif X_train['Sunshine'].loc[index]==0:
            test2.append('Clouds')
        else:
            pass
    else:
        pass
```

```
In [28]: print("There are " + str(test2.count('Sunshine')) + " records of sunshine d
print("There are " + str(test2.count('Clouds')) + " records of cloud data w
print("There are " + str(test2.count('Neither')) + " records of no sunshine
print("There are " + str(test2.count('Both')) + " records of sunshine and c
```

There are 7701 records of sunshine data with 0 cloud coverage.  
 There are 1712 records of cloud data with 0 sunshine hours.  
 There are 4 records of no sunshine or cloud coverage.  
 There are 97226 records of sunshine and cloud coverage.

```
In [29]: test2_data = {'Sunshine with 0 Clouds': 10350, 'Clouds with 0 sunshine': 23
```

```
In [30]: X_train['RainToday'].value_counts(normalize=True)
```

```
Out[30]: No      0.775687
Yes       0.224313
Name: RainToday, dtype: float64
```

```
In [31]: rain_data_today = {'Did Not Rain':113580,'Rained':31880}
```

In [32]:

```

# Create subplot for all of our tests/findings

fig, (ax1, ax2, ax3, ax4, ax5, ax6, ax7, ax8) = plt.subplots(8, 1, figsize=
fig.suptitle('Plots of Cloud data and Sunshine Data')
fig.tight_layout(pad=5.0)

# This is a distribution for sunshine data that eliminates cloud nulls.
ax1.set_title('Chart 1: Dist. of Sunshine Data w/o cloud nulls')
ax1.hist(rain_data_cloud_dropna['Sunshine'])
ax1.set_xlabel('Hours of Sunshine')
ax1.set_ylabel('Sunshine Data')

# This is a distribution for cloud data that eliminates sunshine nulls. It
# hard to impute these values without further biforcation.
ax2.set_title('Chart 2: Dist. of 9am Cloud Data w/o Sunshine nulls')
ax2.hist(rain_data_sunshine_dropna['Cloud9am'])
ax2.set_xlabel('9am Cloud Coverage (0-8)')
ax2.set_ylabel('Cloud Data')

# This is a distribution for cloud data that eliminates sunshine nulls. It
# hard to impute these values without further biforcation.
ax3.set_title('Chart 3: Dist. of 3pm Cloud Data w/o Sunshine nulls')
ax3.hist(rain_data_sunshine_dropna['Cloud3pm'])
ax3.set_xlabel('3pm Cloud Coverage (0-8)')
ax3.set_ylabel('Cloud Data')

# Here we plot the population of nulls and their relationships. Perhaps the
# clouds and sunshine that we can impute. When both records are null, this
ax4.set_title('Chart 4: # of non-nulls by Category (Sunshine vs Cloud Data)')
ax4.bar(test_data.keys(),test_data.values())
ax4.set_xlabel('Type')

# Checking the data for zeros.
ax5.set_title('Chart 5: Non-Zero Data by Category')
ax5.bar(test2_data.keys(),test2_data.values())
ax5.set_xlabel('Type')








# Here we find the distribution for days it does rain, and the relevant mea
ax6.set_title('Chart 6: Sunshine Data On Day it Rains')
ax6.hist(sunshine_when_rain['Sunshine'],bins=14)

# Here we find the distribution for days it doesn't rain, and the relevant
ax7.set_title("Chart 7: Sunshine Data On Day it Doesn't Rain")
ax7.hist(sunshine_no_rain['Sunshine'])


# Here is the distribution for evaporation when sunshine is greater than ze
ax8.set_title("Chart 8: Evaporation Data when Sunshine > 0")
ax8.hist(evaporation_test,bins=100)
plt.savefig('Visualizations/Subplot1.png', bbox_inches = 'tight')
plt.show()

```

Plots of Cloud data and Sunshine Data







Here are the actionable insights gained from the charts above:

- **Chart 1:** For Sunshine Data overall and with cloud nulls filtered out, it appears to be normally distributed with a slight skew to the left. There are also many zeroes, which can be explained by rainy days.
- **Chart 2 and 3:** For Cloud Data overall and with Sunshine nulls removed, the data appears to be hump shaped, with one large hump at a 1 on the scale, and another large hump at a 7 on the scale.
- **Chart 4:** This chart is designed to look at the overall null population, to further understand what data we will need to impute. For our cloud data with sunshine nulls, I would recommend looking further into the rain data to determine how we should handle imputing sunshine values. For sunshine data with cloud nulls, I will impute based off of both the humps; if it rained, I will use the 7 hump, if it did not rain, I will use the 1 hump. If both are null, I will have to use only rain data to determine how I would like to handle these values.
- **Chart 5:** I made this chart to see situations where there are clouds with 0 sunshine, and sunshine with 0 clouds. These values appear to be minimal.
- **Chart 6 and 7:** I created charts here to see how many hours of sunshine there are when it rains, and when it doesn't rain. This could help us biforcate our population of sunshine data so that I am not imputing the mean onto data that doesn't represent the mean.
- **Chart 8:** This shows evaporation data when Sunshine > 0. We can impute the mean here onto our evaporation nulls.

We need to dive deeper into cloud data in order to figure out how to replace nulls. Please see graphs below for more details about our cloud data:

In [33]:

```
# Now I am going to make a subplot looking into the results of our cloud te

fig, ((ax1, ax4), (ax2, ax5), (ax3, ax6)) = plt.subplots(nrows=3, ncols=2,
fig.suptitle('Plots of Cloud data and Sunshine Data')
fig.tight_layout(pad=5.0)

# This is a graph for our cloud data when there is low humidity and it does
ax1.set_title('Chart 9: 9am Cloud Data- Low Humidity, No Rain')
ax1.hist(cloud9_no_rain_lower_humidity)

# This is a graph for our cloud data when there is high humidity and it doe
ax2.set_title('Chart 10: 9am Cloud Data- High Humidity, No Rain')
ax2.hist(cloud9_no_rain_higher_humidity)

# This is a graph for our cloud data when it does rain.
ax3.set_title('Chart 11: 9am Cloud Data- Rains')
ax3.hist(cloud9_rain)

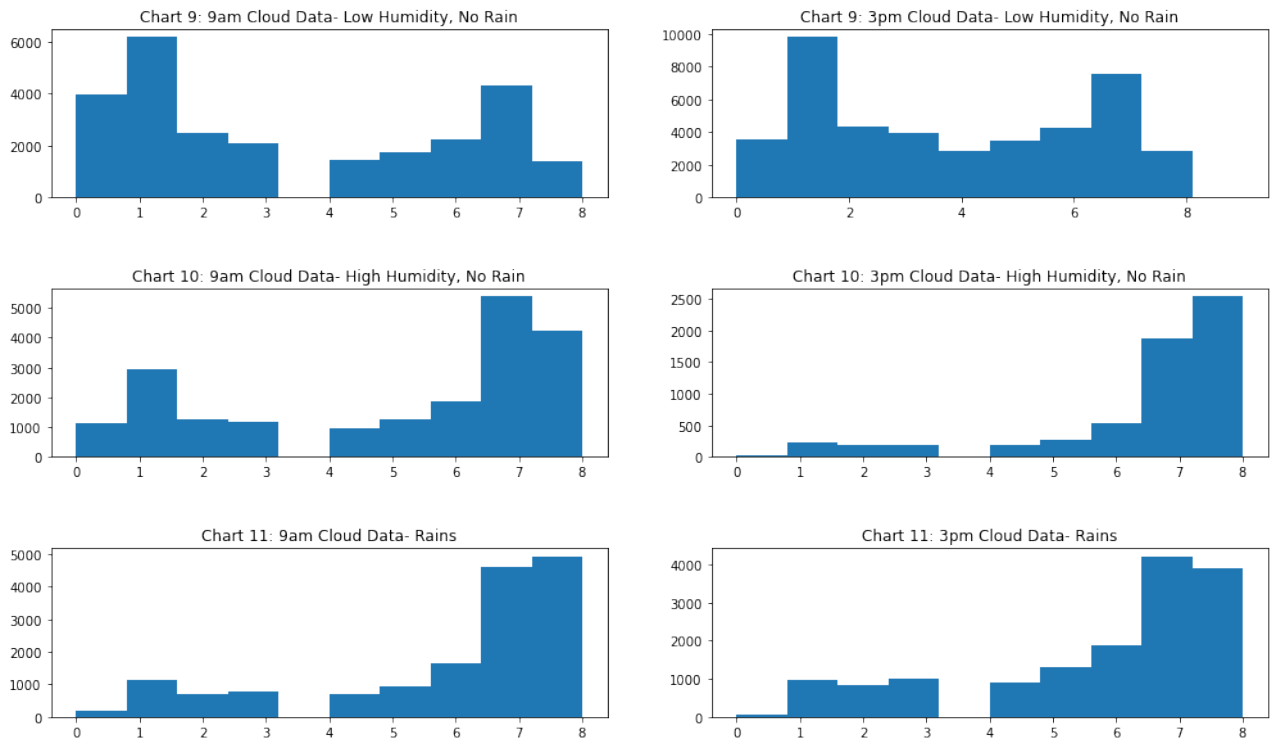
# This is a graph for our cloud data when there is low humidity and it does
ax4.set_title('Chart 9: 3pm Cloud Data- Low Humidity, No Rain')
ax4.hist(cloud3_no_rain_lower_humidity)

# This is a graph for our cloud data when there is high humidity and it doe
ax5.set_title('Chart 10: 3pm Cloud Data- High Humidity, No Rain')
ax5.hist(cloud3_no_rain_higher_humidity)

# This is a graph for our cloud data when it does rain.
ax6.set_title('Chart 11: 3pm Cloud Data- Rains')
ax6.hist(cloud3_rain)

plt.savefig('Visualizations/Subplot2.png', bbox_inches = 'tight')
```

Plots of Cloud data and Sunshine Data



Let's breakdown our Cloud data:

- **Chart 9:** As we can see here, in our data when it doesn't rain and humidity is less than 70%, our Cloud data has a mean closer to 1. I will fill the nulls meeting this criteria using a lower cloud coverage.
- **Chart 10:** As we can see here, in our data when it doesn't rain and humidity is greater than 70%, our Cloud data has a mean closer to 7 or 8. I will fill the nulls meeting this criteria using a higher cloud coverage.
- **Chart 11:** As we can see here, in our data when it does rain, our Cloud data has a mean closer to 7 or 8. I will fill the nulls meeting this criteria using a higher cloud coverage.

Now that I have an understanding of the data, I will perform the train/test split, and begin imputing these null values.

## Section 3: Data Preparation

Before imputing my data, I am going to create numeric columns for our categorical fields using OneHotEncoding.

In [34]:

```
X_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 106644 entries, 0 to 106643
Data columns (total 22 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                   106644 non-null object
1   Location               106644 non-null object
2   MinTemp                106150 non-null float64
3   MaxTemp                106387 non-null float64
4   Rainfall               105540 non-null float64
5   Evaporation            60909 non-null float64
6   Sunshine               55677 non-null float64
7   WindGustDir            99665 non-null object
8   WindGustSpeed          99715 non-null float64
9   WindDir9am             99120 non-null object
10  WindDir3pm             103822 non-null object
11  WindSpeed9am           105646 non-null float64
12  WindSpeed3pm           104681 non-null float64
13  Humidity9am            105326 non-null float64
14  Humidity3pm            103917 non-null float64
15  Pressure9am            96098 non-null float64
16  Pressure3pm            96123 non-null float64
17  Cloud9am               66285 non-null float64
18  Cloud3pm               63757 non-null float64
19  Temp9am                105967 non-null float64
20  Temp3pm                104579 non-null float64
21  RainToday              105540 non-null object
dtypes: float64(16), object(6)
memory usage: 17.9+ MB
```

In [35]:

```
X_train = X_train.drop(['Date'],axis=1)
```

In [36]:

```
# Impute sunshine data based off of our findings from above

X_train['Sunshine'] = np.where(((X_train['Sunshine'].isna()) & (X_train['Ra
X_train['Sunshine'] = np.where(((X_train['Sunshine'].isna()) & (X_train['Ra
```

In [37]:

```
X_train['Sunshine'].isna().any()
```

Out[37]: True

In [38]:

```
X_train['Sunshine'].value_counts()
```

```
Out[38]: 0.0      1716
          10.7     815
```

11.0	812
10.5	779
10.8	771
10.3	750
10.2	732
10.9	732
9.8	726
11.1	708
10.6	708
10.4	705
10.0	704
10.1	704
9.2	698
11.2	686
9.9	667
9.5	647
9.4	639
9.6	629
9.7	617
9.3	605
9.0	565
8.8	558
11.3	557
9.1	555
8.4	548
8.9	526
8.7	525
8.2	522
11.4	515
8.0	505
11.6	491
8.3	480
8.6	468
7.2	463
8.5	462
7.8	457
13.0	448
8.1	448
12.0	446
11.7	431
13.1	425
7.3	423
13.2	421
11.5	420
11.9	418
7.7	418
12.7	418
7.6	417
11.8	414
6.8	411
7.1	408
6.9	406
7.5	403

7.4	401
0.1	399
6.3	398
6.1	396
12.2	394
7.0	388
0.2	388
12.5	383
6.6	380
7.9	370
6.0	366
6.2	366
12.4	365
6.5	364
12.6	363
12.3	359
6.4	356
5.7	356
5.8	353
12.1	353
6.7	343
12.8	342
5.0	340
5.5	339
5.2	337
5.6	325
0.3	325
13.3	319
5.9	317
12.9	316
5.1	310
5.4	310
4.8	307
3.8	301
5.3	300
4.3	295
3.0	295
4.4	293
3.2	292
4.5	291
4.0	291
3.9	291
4.7	290
4.1	285
4.9	279
4.6	273
0.7	263
3.6	262
2.7	255
0.4	253
2.3	252
4.2	251
2.8	245

1.0	244
3.5	239
1.6	236
2.0	235
2.2	234
2.4	233
1.2	233
2.1	230
0.9	229
3.3	229
13.4	227
0.5	226
0.6	226
2.9	224
1.7	224
0.8	223
3.7	220
1.5	217
1.9	216
2.6	215
1.3	213
1.4	213
3.1	212
1.8	211
2.5	210
3.4	205
1.1	201
13.5	147
13.6	126
13.7	88
13.8	48
13.9	14
14.0	10
14.1	5
14.3	4
14.2	2
14.5	1

Name: Sunshine, dtype: int64

```
In [39]: # Impute zero evaporation when sunshine is zero

X_train['Evaporation'] = np.where(((X_train['Evaporation'].isna()) & (X_tra
```

```
In [40]: evaporation_test.sum()/len(evaporation_test)
```

```
Out[40]: 5.037441856155372
```

```
In [41]: # Impute mean evaporation when sunshine is not zero

X_train['Evaporation'] = np.where(((X_train['Evaporation'].isna()) & (X_tra
```

```
In [42]: X_train['Evaporation'].isna().any()
```

```
Out[42]: True
```

```
In [43]: X_train['Evaporation'].value_counts().head()
```

```
Out[43]: 5.0    4186
         4.0    2494
         8.0    1938
         2.2    1567
         2.0    1494
         Name: Evaporation, dtype: int64
```

```
In [44]: print(cloud9_no_rain_lower_humidity.value_counts())
         print(cloud9_no_rain_higher_humidity.value_counts())
         print(cloud9_rain.value_counts())
         print(cloud3_no_rain_lower_humidity.value_counts())
         print(cloud3_no_rain_higher_humidity.value_counts())
         print(cloud3_rain.value_counts())
```

```
1.0    6175
7.0    4297
0.0    3953
2.0    2453
6.0    2237
3.0    2092
5.0    1743
4.0    1422
8.0    1391
         Name: Cloud9am, dtype: int64
7.0    5366
8.0    4206
1.0    2952
6.0    1884
5.0    1282
2.0    1261
3.0    1179
0.0    1132
4.0     949
         Name: Cloud9am, dtype: int64
8.0    4938
7.0    4617
6.0    1631
1.0    1132
```



```

5.0      955
3.0      794
4.0      710
2.0      691
0.0      174
Name: Cloud9am, dtype: int64
1.0      9839
7.0      7523
2.0      4320
6.0      4268
3.0      3943
0.0      3560
5.0      3474
8.0      2850
4.0      2834
9.0         1
Name: Cloud3pm, dtype: int64
8.0      2529
7.0      1878
6.0       540
5.0       282
1.0       223
4.0       199
3.0       187
2.0       184
0.0        31
Name: Cloud3pm, dtype: int64
7.0      4215
8.0      3907
6.0      1865
5.0      1302
3.0       988
1.0       967
4.0       915
2.0       827
0.0        69
Name: Cloud3pm, dtype: int64

```

In [45]:

```

# Imputing cloud data based on our findings from section 2

X_train['Cloud9am'] = np.where(((X_train['Cloud9am']).isna()) & (X_train['Ra
X_train['Cloud3pm'] = np.where(((X_train['Cloud3pm']).isna()) & (X_train['Ra
X_train['Cloud9am'] = np.where(((X_train['Cloud9am']).isna()) & (X_train['Ra
X_train['Cloud9am'] = np.where(((X_train['Cloud9am']).isna()) & (X_train['Ra
X_train['Cloud3pm'] = np.where(((X_train['Cloud3pm']).isna()) & (X_train['Ra
X_train['Cloud3pm'] = np.where(((X_train['Cloud3pm']).isna()) & (X_train['Ra

```

In [46]:

```
X_train['Cloud9am'].value_counts()
```

```
Out[46]: 7.0    14805
         1.0    11618
         8.0    10767
         0.0     6427
         6.0     5999
         2.0     4847
         3.0     4385
         5.0     4177
         4.0     3260
         Name: Cloud9am, dtype: int64
```

```
In [47]: X_train['Cloud3pm'].value_counts()
```

```
Out[47]: 7.0    13512
         1.0    11081
         8.0     9328
         6.0     6642
         2.0     5346
         3.0     5126
         5.0     5049
         4.0     3952
         0.0     3720
         9.0         1
         Name: Cloud3pm, dtype: int64
```

```
In [48]: # Imputing remaining columns based on mean, from our charts in section 2

         imputer = SimpleImputer(strategy='most_frequent')
         imputer = imputer.fit(X_train)
         X_train.iloc[:, :] = imputer.transform(X_train)
```

```
In [49]: categorical_data = X_train[['Location', 'WindGustDir', 'WindDir9am', 'WindDir3pm']]

         ohe = OneHotEncoder()

         # Fit the dummy variables to an array
         X = ohe.fit_transform(categorical_data.values).toarray()
         y = ohe.get_feature_names()

         # To add this back into the original dataframe
         dfOneHot = pd.DataFrame(X, columns = y)
         X_train = pd.concat([X_train, dfOneHot], axis=1)

         # Dropping the base columns, and to avoid multicollinearity, dropping one of them
         X_train = X_train.drop(['Location', 'WindGustDir', 'WindDir9am', 'WindDir3pm'],
                                axis=1)

         # Printing to verify
         print(X_train.head())
```

```
MinTemp  MaxTemp  Rainfall  Evaporation  Sunshine  WindGustSpeed  \
```

0	14.9	30.3	0.0	7.4	10.9	33.0
1	14.6	21.5	0.2	5.0	0.0	46.0
2	9.0	23.7	0.0	5.0	0.0	28.0
3	15.3	24.0	0.0	8.2	12.1	35.0
4	17.3	37.5	0.0	8.6	11.4	39.0

	WindSpeed9am	WindSpeed3pm	Humidity9am	Humidity3pm	Pressure9am	\
0	15.0	11.0	19.0	12.0	1021.1	
1	26.0	28.0	65.0	57.0	1012.6	
2	11.0	15.0	59.0	45.0	1017.9	
3	4.0	15.0	63.0	82.0	1018.1	
4	9.0	15.0	26.0	12.0	1009.6	

	Pressure3pm	Cloud9am	Cloud3pm	Temp9am	Temp3pm	x0_Adelaide	x0_Albury
0	1017.8	1.0	1.0	22.2	29.7	0.0	0.0
1	1013.5	7.0	7.0	17.5	18.6	0.0	0.0
2	1015.1	7.0	7.0	14.6	23.1	0.0	0.0
3	1016.7	3.0	3.0	21.8	21.8	0.0	0.0
4	1006.2	7.0	4.0	23.8	35.7	0.0	0.0

	x0_AliceSprings	x0_BadgerysCreek	x0_Ballarat	x0_Bendigo	x0_Brisbane
0	0.0	0.0	0.0	0.0	0.0
1	0.0	0.0	0.0	0.0	0.0
2	0.0	0.0	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0	0.0
4	0.0	0.0	0.0	0.0	0.0

	x0_Cairns	x0_Canberra	x0_Cobar	x0_CoffsHarbour	x0_Dartmoor	x0_Darwin
0	0.0	0.0	0.0	0.0	0.0	0.0
1	0.0	0.0	0.0	0.0	0.0	0.0
2	0.0	0.0	0.0	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0	0.0	0.0
4	0.0	0.0	0.0	0.0	0.0	0.0

	x0_GoldCoast	x0_Hobart	x0_Katherine	x0_Launceston	x0_Melbourne	\
0	0.0	0.0	0.0	0.0	0.0	
1	0.0	0.0	0.0	0.0	0.0	
2	0.0	0.0	0.0	0.0	0.0	
3	0.0	0.0	0.0	0.0	0.0	
4	0.0	0.0	0.0	0.0	0.0	

	x0_MelbourneAirport	x0_Mildura	x0_Moree	x0_MountGambier	x0_MountGinin
0	0.0	0.0	0.0	0.0	0.
1	0.0	0.0	0.0	0.0	0.
2	0.0	0.0	0.0	0.0	0.
3	0.0	0.0	0.0	0.0	0.

```

0
4          0.0          1.0          0.0          0.0          0.
0

    x0_Newcastle  x0_Nhil  x0_NorahHead  x0_NorfolkIsland  x0_Nuriootpa  \
0          0.0      0.0          0.0          0.0          0.0
1          0.0      0.0          0.0          0.0          0.0
2          0.0      0.0          0.0          0.0          0.0
3          0.0      0.0          0.0          0.0          0.0
4          0.0      0.0          0.0          0.0          0.0

    x0_PearceRAAF  x0_Perth  x0_PerthAirport  x0_Portland  x0_Richmond  \
0          0.0      0.0          0.0          0.0          0.0
1          0.0      0.0          0.0          0.0          0.0
2          0.0      0.0          0.0          0.0          0.0
3          0.0      0.0          0.0          0.0          0.0
4          0.0      0.0          0.0          0.0          0.0

    x0_Sale  x0_SalmonGums  x0_Sydney  x0_SydneyAirport  x0_Townsville  \
0          0.0          0.0          0.0          0.0          0.0
1          0.0          0.0          0.0          0.0          0.0
2          0.0          1.0          0.0          0.0          0.0
3          0.0          0.0          0.0          0.0          0.0
4          0.0          0.0          0.0          0.0          0.0

    x0_Tuggeranong  x0_Uluru  x0_WaggaWagga  x0_Walpole  x0_Watsonia  \
0          0.0      0.0          0.0          0.0          0.0
1          0.0      0.0          0.0          0.0          0.0
2          0.0      0.0          0.0          0.0          0.0
3          0.0      0.0          0.0          0.0          0.0
4          0.0      0.0          0.0          0.0          0.0

    x0_Williamtown  x0_Witchcliffe  x0_Wollongong  x0_Woomera  x1_E  x1_ENE
\
0          0.0          0.0          0.0          0.0          1.0  0.0  0.0
1          0.0          1.0          0.0          0.0          0.0  0.0  0.0
2          0.0          0.0          0.0          0.0          0.0  0.0  0.0
3          0.0          0.0          0.0          0.0          0.0  0.0  0.0
4          0.0          0.0          0.0          0.0          0.0  0.0  0.0

    x1_ESE  x1_N  x1_NE  x1_NNW  x1_NW  x1_S  x1_SE  x1_SSE  x1_SSW  x1_SW  \
0          0.0  0.0  0.0  0.0  0.0  1.0  0.0  0.0  0.0  0.0
1          0.0  0.0  0.0  0.0  0.0  0.0  0.0  1.0  0.0  0.0
2          0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
3          0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
4          0.0  1.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0

    x1_W  x1_WNW  x1_WSW  x2_E  x2_ENE  x2_N  x2_NE  x2_NNE  x2_NNW  x2_NW  \
0          0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
1          0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
2          1.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
3          1.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
4          0.0  0.0  0.0  0.0  0.0  0.0  1.0  0.0  0.0  0.0

```

	x2_S	x2_SE	x2_SSE	x2_SSW	x2_SW	x2_W	x2_WNW	x2_WSW	x3_E	x3_ESE	\
0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
1	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
2	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	
3	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	
4	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	

	x3_N	x3_NE	x3_NNE	x3_NNW	x3_NW	x3_S	x3_SE	x3_SSE	x3_SSW	x3_SW	\
0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	
1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	
2	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
3	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	
4	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	

	x3_W	x3_WNW	x3_WSW	x4_Yes
0	0.0	0.0	0.0	0.0
1	0.0	0.0	0.0	0.0
2	1.0	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0
4	0.0	0.0	0.0	0.0

```
In [50]: # Imputing remaining columns based on mean, from our charts in section 2

imputer = SimpleImputer(strategy='most_frequent', missing_values=np.nan)
imputer = imputer.fit(X_train)
X_train.iloc[:, :] = imputer.transform(X_train)
```

```
In [51]: X_train.info(1, null_counts=True)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 106644 entries, 0 to 106643
Data columns (total 109 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   MinTemp                             106644 non-null float64
1   MaxTemp                             106644 non-null float64
2   Rainfall                            106644 non-null float64
3   Evaporation                         106644 non-null float64
4   Sunshine                            106644 non-null float64
5   WindGustSpeed                       106644 non-null float64
6   WindSpeed9am                       106644 non-null float64
7   WindSpeed3pm                       106644 non-null float64
8   Humidity9am                        106644 non-null float64
9   Humidity3pm                        106644 non-null float64
10  Pressure9am                         106644 non-null float64
11  Pressure3pm                         106644 non-null float64
12  Cloud9am                            106644 non-null float64
13  Cloud3pm                            106644 non-null float64
14  Temp9am                             106644 non-null float64
```

15	Temp3pm	106644	non-null	float64
16	x0_Adelaide	106644	non-null	float64
17	x0_Albury	106644	non-null	float64
18	x0_AliceSprings	106644	non-null	float64
19	x0_BadgerysCreek	106644	non-null	float64
20	x0_Ballarat	106644	non-null	float64
21	x0_Bendigo	106644	non-null	float64
22	x0_Brisbane	106644	non-null	float64
23	x0_Cairns	106644	non-null	float64
24	x0_Canberra	106644	non-null	float64
25	x0_Cobar	106644	non-null	float64
26	x0_CoffsHarbour	106644	non-null	float64
27	x0_Dartmoor	106644	non-null	float64
28	x0_Darwin	106644	non-null	float64
29	x0_GoldCoast	106644	non-null	float64
30	x0_Hobart	106644	non-null	float64
31	x0_Katherine	106644	non-null	float64
32	x0_Launceston	106644	non-null	float64
33	x0_Melbourne	106644	non-null	float64
34	x0_MelbourneAirport	106644	non-null	float64
35	x0_Mildura	106644	non-null	float64
36	x0_Moree	106644	non-null	float64
37	x0_MountGambier	106644	non-null	float64
38	x0_MountGinini	106644	non-null	float64
39	x0_Newcastle	106644	non-null	float64
40	x0_Nhil	106644	non-null	float64
41	x0_NorahHead	106644	non-null	float64
42	x0_NorfolkIsland	106644	non-null	float64
43	x0_Nuriootpa	106644	non-null	float64
44	x0_PearceRAAF	106644	non-null	float64
45	x0_Perth	106644	non-null	float64
46	x0_PerthAirport	106644	non-null	float64
47	x0_Portland	106644	non-null	float64
48	x0_Richmond	106644	non-null	float64
49	x0_Sale	106644	non-null	float64
50	x0_SalmonGums	106644	non-null	float64
51	x0_Sydney	106644	non-null	float64
52	x0_SydneyAirport	106644	non-null	float64
53	x0_Townsville	106644	non-null	float64
54	x0_Tuggeranong	106644	non-null	float64
55	x0_Uluru	106644	non-null	float64
56	x0_WaggaWagga	106644	non-null	float64
57	x0_Walpole	106644	non-null	float64
58	x0_Watsonia	106644	non-null	float64
59	x0_Williamtown	106644	non-null	float64
60	x0_Witchcliffe	106644	non-null	float64
61	x0_Wollongong	106644	non-null	float64
62	x0_Woomera	106644	non-null	float64
63	x1_E	106644	non-null	float64
64	x1_ENE	106644	non-null	float64
65	x1_ESE	106644	non-null	float64
66	x1_N	106644	non-null	float64
67	x1_NE	106644	non-null	float64

```
68  x1_NNW      106644 non-null float64
69  x1_NW       106644 non-null float64
70  x1_S        106644 non-null float64
71  x1_SE       106644 non-null float64
72  x1_SSE      106644 non-null float64
73  x1_SSW      106644 non-null float64
74  x1_SW       106644 non-null float64
75  x1_W        106644 non-null float64
76  x1_WNW      106644 non-null float64
77  x1_WSW      106644 non-null float64
78  x2_E        106644 non-null float64
79  x2_ENE      106644 non-null float64
80  x2_N        106644 non-null float64
81  x2_NE       106644 non-null float64
82  x2_NNE      106644 non-null float64
83  x2_NNW      106644 non-null float64
84  x2_NW       106644 non-null float64
85  x2_S        106644 non-null float64
86  x2_SE       106644 non-null float64
87  x2_SSE      106644 non-null float64
88  x2_SSW      106644 non-null float64
89  x2_SW       106644 non-null float64
90  x2_W        106644 non-null float64
91  x2_WNW      106644 non-null float64
92  x2_WSW      106644 non-null float64
93  x3_E        106644 non-null float64
94  x3_ESE      106644 non-null float64
95  x3_N        106644 non-null float64
96  x3_NE       106644 non-null float64
97  x3_NNE      106644 non-null float64
98  x3_NNW      106644 non-null float64
99  x3_NW       106644 non-null float64
100 x3_S        106644 non-null float64
101 x3_SE       106644 non-null float64
102 x3_SSE      106644 non-null float64
103 x3_SSW      106644 non-null float64
104 x3_SW       106644 non-null float64
105 x3_W        106644 non-null float64
106 x3_WNW      106644 non-null float64
107 x3_WSW      106644 non-null float64
108 x4_Yes      106644 non-null float64
dtypes: float64(109)
memory usage: 88.7 MB
```

```
In [52]: y_train.count()
```

```
Out[52]: 106644
```

```
In [53]: X_test.info(1, null_counts=True)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 35549 entries, 0 to 35548
Data columns (total 22 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                  35549 non-null  object
1   Location              35549 non-null  object
2   MinTemp               35406 non-null  float64
3   MaxTemp               35484 non-null  float64
4   Rainfall              35247 non-null  float64
5   Evaporation           20441 non-null  float64
6   Sunshine              18700 non-null  float64
7   WindGustDir           33198 non-null  object
8   WindGustSpeed         33208 non-null  float64
9   WindDir9am            33060 non-null  object
10  WindDir3pm            34593 non-null  object
11  WindSpeed9am          35199 non-null  float64
12  WindSpeed3pm          34882 non-null  float64
13  Humidity9am           35093 non-null  float64
14  Humidity3pm           34666 non-null  float64
15  Pressure9am           32081 non-null  float64
16  Pressure3pm           32089 non-null  float64
17  Cloud9am              22251 non-null  float64
18  Cloud3pm              21342 non-null  float64
19  Temp9am               35322 non-null  float64
20  Temp3pm               34888 non-null  float64
21  RainToday             35247 non-null  object
dtypes: float64(16), object(6)
memory usage: 6.0+ MB
```

```
In [54]: X_test.reset_index(inplace=True)
```

```
In [55]: X_test = X_test.drop(columns=['index'],axis=1)
```

```
In [56]: # Performing same imputations as train data

X_test['Sunshine'] = np.where(((X_test['Sunshine'].isna()) & (X_test['RainT
X_test['Sunshine'] = np.where(((X_test['Sunshine'].isna()) & (X_test['RainT
X_test['Evaporation'] = np.where(((X_test['Evaporation'].isna()) & (X_test[
X_test['Evaporation'] = np.where(((X_test['Evaporation'].isna()) & (X_test[
X_test['Cloud9am'] = np.where(((X_test['Cloud9am'].isna()) & (X_test['RainT
X_test['Cloud3pm'] = np.where(((X_test['Cloud3pm'].isna()) & (X_test['RainT
X_test['Cloud9am'] = np.where(((X_test['Cloud9am'].isna()) & (X_test['RainT
X_test['Cloud9am'] = np.where(((X_test['Cloud9am'].isna()) & (X_test['RainT
X_test['Cloud3pm'] = np.where(((X_test['Cloud3pm'].isna()) & (X_test['RainT
X_test['Cloud3pm'] = np.where(((X_test['Cloud3pm'].isna()) & (X_test['RainT
```



In [57]:

```
# Performing same imputations as train data

imputer = SimpleImputer(strategy='most_frequent')
imputer = imputer.fit(X_test)
X_test.iloc[:, :] = imputer.transform(X_test)
```

In [58]:

```
categorical_data = X_test[['Location', 'WindGustDir', 'WindDir9am', 'WindDir3p

ohe = OneHotEncoder()

# Fit the dummy variables to an array
X = ohe.fit_transform(categorical_data.values).toarray()
y = ohe.get_feature_names()

# To add this back into the original dataframe
dfOneHot = pd.DataFrame(X, columns = y)
X_test = pd.concat([X_test, dfOneHot], axis=1)

# Dropping the country column
X_test = X_test.drop(['Date', 'Location', 'WindGustDir', 'WindDir9am', 'WindDir

# Printing to verify
print(X_test.head())
```

	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustSpeed	\
0	7.1	13.0	8.8	5.0	0.0	41.0	
1	13.2	18.3	0.0	5.0	0.0	48.0	
2	9.2	22.7	0.0	5.0	11.1	52.0	
3	15.3	26.1	0.0	10.4	0.0	44.0	
4	11.9	31.8	0.0	5.0	4.1	72.0	

	WindSpeed9am	WindSpeed3pm	Humidity9am	Humidity3pm	Pressure9am	\
0	24.0	22.0	100.0	98.0	1001.7	
1	24.0	20.0	73.0	73.0	1027.6	
2	26.0	20.0	45.0	25.0	1030.1	
3	24.0	19.0	48.0	40.0	1013.2	
4	6.0	19.0	89.0	25.0	1006.7	

	Pressure3pm	Cloud9am	Cloud3pm	Temp9am	Temp3pm	x0_Adelaide	x0_Albury
0	1005.4	8.0	8.0	8.6	11.5	0.0	0.0
1	1023.8	7.0	7.0	14.2	17.0	0.0	0.0
2	1025.9	0.0	0.0	15.1	22.5	0.0	0.0
3	1009.8	7.0	7.0	17.5	24.3	0.0	0.0
4	1001.0	7.0	6.0	16.2	27.4	0.0	0.0

	x0_AliceSprings	x0_BadgerysCreek	x0_Ballarat	x0_Bendigo	x0_Brisbane
0	0.0	0.0	1.0	0.0	0.0
1	0.0	0.0	0.0	0.0	0.0

2	0.0	0.0	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0	0.0
4	0.0	0.0	0.0	0.0	0.0

	x0_Cairns	x0_Canberra	x0_Cobar	x0_CoffsHarbour	x0_Dartmoor	x0_Darwin
\						
0	0.0	0.0	0.0	0.0	0.0	0.0
1	0.0	0.0	0.0	0.0	0.0	0.0
2	0.0	0.0	0.0	0.0	0.0	0.0
3	0.0	0.0	1.0	0.0	0.0	0.0
4	0.0	0.0	0.0	0.0	0.0	0.0

	x0_GoldCoast	x0_Hobart	x0_Katherine	x0_Launceston	x0_Melbourne	\
0	0.0	0.0	0.0	0.0	0.0	
1	0.0	0.0	0.0	0.0	0.0	
2	0.0	0.0	0.0	0.0	0.0	
3	0.0	0.0	0.0	0.0	0.0	
4	0.0	0.0	0.0	0.0	0.0	

	x0_MelbourneAirport	x0_Mildura	x0_Moree	x0_MountGambier	x0_MountGinin	i \
0		0.0	0.0	0.0	0.0	0.
0						
1		0.0	0.0	0.0	0.0	0.
0						
2		0.0	0.0	0.0	0.0	0.
0						
3		0.0	0.0	0.0	0.0	0.
0						
4		0.0	0.0	0.0	0.0	0.
0						

	x0_Newcastle	x0_Nhil	x0_NorahHead	x0_NorfolkIsland	x0_Nuriootpa	\
0	0.0	0.0	0.0	0.0	0.0	
1	0.0	0.0	0.0	0.0	0.0	
2	0.0	0.0	0.0	0.0	0.0	
3	0.0	0.0	0.0	0.0	0.0	
4	0.0	0.0	0.0	0.0	0.0	

	x0_PearceRAAF	x0_Perth	x0_PerthAirport	x0_Portland	x0_Richmond	\
0	0.0	0.0	0.0	0.0	0.0	
1	0.0	0.0	0.0	0.0	0.0	
2	0.0	0.0	1.0	0.0	0.0	
3	0.0	0.0	0.0	0.0	0.0	
4	0.0	0.0	0.0	0.0	0.0	

	x0_Sale	x0_SalmonGums	x0_Sydney	x0_SydneyAirport	x0_Townsville	\
0	0.0	0.0	0.0	0.0	0.0	
1	0.0	0.0	0.0	0.0	0.0	
2	0.0	0.0	0.0	0.0	0.0	
3	0.0	0.0	0.0	0.0	0.0	
4	1.0	0.0	0.0	0.0	0.0	

	x0_Tuggeranong	x0_Uluru	x0_WaggaWagga	x0_Walpole	x0_Watsonia	\
0	0.0	0.0	0.0	0.0	0.0	
1	0.0	0.0	0.0	1.0	0.0	
2	0.0	0.0	0.0	0.0	0.0	
3	0.0	0.0	0.0	0.0	0.0	
4	0.0	0.0	0.0	0.0	0.0	

	x0_Williamtown	x0_Witchcliffe	x0_Wollongong	x0_Woomera	x1_E	x1_ENE
0	0.0	0.0	0.0	0.0	0.0	0.0
1	0.0	0.0	0.0	0.0	1.0	0.0
2	0.0	0.0	0.0	0.0	0.0	1.0
3	0.0	0.0	0.0	0.0	1.0	0.0
4	0.0	0.0	0.0	0.0	0.0	0.0

	x1_ESE	x1_N	x1_NE	x1_NNW	x1_NW	x1_S	x1_SE	x1_SSE	x1_SSW	x1_SW	\
0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
2	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
3	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
4	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	

	x1_W	x1_WNW	x1_WSW	x2_E	x2_ENE	x2_N	x2_NE	x2_NNE	x2_NNW	x2_NW	\
0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	
1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
2	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	
3	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	
4	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	

	x2_S	x2_SE	x2_SSE	x2_SSW	x2_SW	x2_W	x2_WNW	x2_WSW	x3_E	x3_ESE	\
0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	
2	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	
3	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
4	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	

	x3_N	x3_NE	x3_NNE	x3_NNW	x3_NW	x3_S	x3_SE	x3_SSE	x3_SSW	x3_SW	\
0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
2	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
3	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
4	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	

	x3_W	x3_WNW	x3_WSW	x4_Yes
0	0.0	1.0	0.0	1.0
1	0.0	0.0	0.0	0.0
2	0.0	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0
4	0.0	0.0	0.0	0.0

```
In [59]: x_test.info(1, null_counts=True)
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 35549 entries, 0 to 35548
```

```
Data columns (total 109 columns):
```

#	Column	Non-Null Count	Dtype
0	MinTemp	35549 non-null	float64
1	MaxTemp	35549 non-null	float64
2	Rainfall	35549 non-null	float64
3	Evaporation	35549 non-null	float64
4	Sunshine	35549 non-null	float64
5	WindGustSpeed	35549 non-null	float64
6	WindSpeed9am	35549 non-null	float64
7	WindSpeed3pm	35549 non-null	float64
8	Humidity9am	35549 non-null	float64
9	Humidity3pm	35549 non-null	float64
10	Pressure9am	35549 non-null	float64
11	Pressure3pm	35549 non-null	float64
12	Cloud9am	35549 non-null	float64
13	Cloud3pm	35549 non-null	float64
14	Temp9am	35549 non-null	float64
15	Temp3pm	35549 non-null	float64
16	x0_Adelaide	35549 non-null	float64
17	x0_Albury	35549 non-null	float64
18	x0_AliceSprings	35549 non-null	float64
19	x0_BadgerysCreek	35549 non-null	float64
20	x0_Ballarat	35549 non-null	float64
21	x0_Bendigo	35549 non-null	float64
22	x0_Brisbane	35549 non-null	float64
23	x0_Cairns	35549 non-null	float64
24	x0_Canberra	35549 non-null	float64
25	x0_Cobar	35549 non-null	float64
26	x0_CoffsHarbour	35549 non-null	float64
27	x0_Dartmoor	35549 non-null	float64
28	x0_Darwin	35549 non-null	float64
29	x0_GoldCoast	35549 non-null	float64
30	x0_Hobart	35549 non-null	float64
31	x0_Katherine	35549 non-null	float64
32	x0_Launceston	35549 non-null	float64
33	x0_Melbourne	35549 non-null	float64
34	x0_MelbourneAirport	35549 non-null	float64
35	x0_Mildura	35549 non-null	float64
36	x0_Moree	35549 non-null	float64
37	x0_MountGambier	35549 non-null	float64
38	x0_MountGinini	35549 non-null	float64
39	x0_Newcastle	35549 non-null	float64
40	x0_Nhil	35549 non-null	float64
41	x0_NorahHead	35549 non-null	float64
42	x0_NorfolkIsland	35549 non-null	float64
43	x0_Nuriootpa	35549 non-null	float64
44	x0_PearceRAAF	35549 non-null	float64
45	x0_Perth	35549 non-null	float64
46	x0_PerthAirport	35549 non-null	float64
47	x0_Portland	35549 non-null	float64

48	x0_Richmond	35549	non-null	float64
49	x0_Sale	35549	non-null	float64
50	x0_SalmonGums	35549	non-null	float64
51	x0_Sydney	35549	non-null	float64
52	x0_SydneyAirport	35549	non-null	float64
53	x0_Townsville	35549	non-null	float64
54	x0_Tuggeranong	35549	non-null	float64
55	x0_Uluru	35549	non-null	float64
56	x0_WaggaWagga	35549	non-null	float64
57	x0_Walpole	35549	non-null	float64
58	x0_Watsonia	35549	non-null	float64
59	x0_Williamtown	35549	non-null	float64
60	x0_Witchcliffe	35549	non-null	float64
61	x0_Wollongong	35549	non-null	float64
62	x0_Woomera	35549	non-null	float64
63	x1_E	35549	non-null	float64
64	x1_ENE	35549	non-null	float64
65	x1_ESE	35549	non-null	float64
66	x1_N	35549	non-null	float64
67	x1_NE	35549	non-null	float64
68	x1_NNW	35549	non-null	float64
69	x1_NW	35549	non-null	float64
70	x1_S	35549	non-null	float64
71	x1_SE	35549	non-null	float64
72	x1_SSE	35549	non-null	float64
73	x1_SSW	35549	non-null	float64
74	x1_SW	35549	non-null	float64
75	x1_W	35549	non-null	float64
76	x1_WNW	35549	non-null	float64
77	x1_WSW	35549	non-null	float64
78	x2_E	35549	non-null	float64
79	x2_ENE	35549	non-null	float64
80	x2_N	35549	non-null	float64
81	x2_NE	35549	non-null	float64
82	x2_NNE	35549	non-null	float64
83	x2_NNW	35549	non-null	float64
84	x2_NW	35549	non-null	float64
85	x2_S	35549	non-null	float64
86	x2_SE	35549	non-null	float64
87	x2_SSE	35549	non-null	float64
88	x2_SSW	35549	non-null	float64
89	x2_SW	35549	non-null	float64
90	x2_W	35549	non-null	float64
91	x2_WNW	35549	non-null	float64
92	x2_WSW	35549	non-null	float64
93	x3_E	35549	non-null	float64
94	x3_ESE	35549	non-null	float64
95	x3_N	35549	non-null	float64
96	x3_NE	35549	non-null	float64
97	x3_NNE	35549	non-null	float64
98	x3_NNW	35549	non-null	float64
99	x3_NW	35549	non-null	float64
100	x3_S	35549	non-null	float64

```
101 x3_SE          35549 non-null float64
102 x3_SSE         35549 non-null float64
103 x3_SSW         35549 non-null float64
104 x3_SW          35549 non-null float64
105 x3_W           35549 non-null float64
106 x3_WNW         35549 non-null float64
107 x3_WSW         35549 non-null float64
108 x4_Yes         35549 non-null float64
dtypes: float64(109)
memory usage: 29.6 MB
```

```
In [60]: y_test.count()
```

```
Out[60]: 35549
```

## Section 4: Modeling

To begin, I will build a logistic model. After that, I will be building a kNN model and a decision tree model, and will pick the best model from the three.

The metric I focus on the most will be recall, because if there are false negatives in my model, then it could rain on citizens who expected a dry day. As the Bureau of Meteorology has been underestimating rain over the past five years, we are trying to minimize false negatives as much as possible.

### Logistic Regression

```
In [61]: # Create a Baseline Model

logreg = LogisticRegression()
model_log = logreg.fit(X_train, y_train)
model_log.score(X_test, y_test)
y_pred_lr = model_log.predict(X_test)
```

```
In [62]: recall_score(y_test, y_pred_lr)
```

```
Out[62]: 0.4690890739338885
```

I will now run various models through a Pipeline to see if I can find a combination of parameters to bring down our performance metric.

```
In [63]: # Create pipeline

lr_pipeline_1 = Pipeline([('ss', StandardScaler()),
                           ('lr', LogisticRegression())])
```

```
In [64]: # Create grid using GridSearchCV

lr_grid = [{'lr__random_state':[42],
            'lr__C': [1e-3, 1e-2, 1, 1e2, 1e3]}]
```

```
In [65]: # Create grid using GridSearchCV

lr_gridsearch = GridSearchCV(estimator=lr_pipeline_1,
                              param_grid=lr_grid,
                              scoring='recall',
                              cv=3)
```

```
In [66]: # Re-distributing our population using SMOTE, to account for accuracy bias

smote = SMOTE()
X_train_resampled, y_train_resampled = smote.fit_resample(X_train, y_train)
```

```
In [67]: y_train_resampled.value_counts()
```

```
Out[67]: 1.0    82693
         0.0    82693
         Name: RainTomorrow, dtype: int64
```

```
In [68]: # Fit the training data
lr_gridsearch.fit(X_train_resampled, y_train_resampled)
print(lr_gridsearch.score(X_test, y_test))

0.7721423164269493
```

```
In [69]: # Print best parameters

print(lr_gridsearch.best_params_)


{'lr__C': 1, 'lr__random_state': 42}
```

In [70]:

```
# Create confusion matrix for baseline model

plot_confusion_matrix(logreg, X_test, y_test,
                      cmap=plt.cm.Blues)

plt.grid(False)
plt.savefig('Visualizations/LogRegBase.png', bbox_inches = 'tight')
plt.show()
```




In [71]:

```
# Create confusion matrix for best model

plot_confusion_matrix(lr_gridsearch, X_test, y_test,
                      cmap=plt.cm.Blues)

plt.grid(False)
plt.savefig('Visualizations/LogRegBest.png', bbox_inches = 'tight')
plt.show()
```





## k-Nearest Neighbors

```
In [72]: # Create baseline model

knn_baseline_model = KNeighborsClassifier()
model_knn = knn_baseline_model.fit(X_train, y_train)
model_knn.score(X_test, y_test)
y_pred_knn = model_knn.predict(X_test)
```

```
In [73]: recall_score(y_test, y_pred_knn)
```

```
Out[73]: 0.5099671965682564
```

Next, I will build a model using Pipeline and kNN to check score using a set of different parameters:

```
In [74]: # Creating 3 different tests using different parameters

knn_pipeline = Pipeline([('ss', StandardScaler()),
                          ('knn', KNeighborsClassifier())])
```

```
In [75]: # Create grid using GridSearchCV

knn_grid = [{'knn__n_neighbors': [2, 5, 7, 10]}]
```

```
In [76]: # Create grid using GridSearchCV

knn_gridsearch = GridSearchCV(estimator=knn_pipeline,
                              param_grid=knn_grid,
                              scoring='recall',
                              cv=3)
```

```
In [77]: # Fit the training data

knn_gridsearch.fit(X_train_resampled, y_train_resampled)
print(knn_gridsearch.score(X_test, y_test))
```

```
0.5547564976028262
```

```
In [78]: print(knn_gridsearch.best_params_)


{'knn__n_neighbors': 5}
```

In [79]:

```
# Create confusion matrix for baseline

plot_confusion_matrix(knn_baseline_model, X_test, y_test,
                      cmap=plt.cm.Blues)

plt.grid(False)
plt.savefig('Visualizations/KNNbaseline.png', bbox_inches = 'tight')
plt.show()
```




In [80]:

```
# Create confusion matrix for best

plot_confusion_matrix(knn_gridsearch, X_test, y_test,
                      cmap=plt.cm.Blues)

plt.grid(False)
plt.savefig('Visualizations/KNNbest.png', bbox_inches = 'tight')
plt.show()
```



As we can see above, our best model seems to cap at a score of around .548. I will now use decision trees and Grid Search CV to see if they generate a stronger model.

## Decision Trees

```
In [81]: # Create a Baseline Model

dt_base = DecisionTreeClassifier()
model_dt = dt_base.fit(X_train, y_train)
model_dt.score(X_test, y_test)
y_pred_dt = model_dt.predict(X_test)
```

```
In [82]: recall_score(y_test, y_pred_dt)
```

```
Out[82]: 0.5282614181175876
```

```
In [83]: # Create pipeline

dtree_pipeline = Pipeline([('ss', StandardScaler()),
                             ('dt', DecisionTreeClassifier())])
```

```
In [84]: # Use GridSearchCV to create different models

dt_grid = [{'dt__criterion': ['gini', 'entropy'],
            'dt__max_leaf_nodes': [5, 10, 20, None],
            'dt__max_features': ['auto', 'sqrt', 'log2', None],
            'dt__random_state': [42]}]
```

```
In [85]: # Use GridSearchCV to create different models

dt_gridsearch = GridSearchCV(estimator=dtree_pipeline,
                              param_grid=dt_grid,
                              scoring='recall',
                              cv=3)
```

```
In [86]: # Fit model and print

dt_gridsearch.fit(X_train_resampled, y_train_resampled)
print(dt_gridsearch.score(X_test, y_test))
```

```
0.7763058289174868
```

```
In [87]: print(dt_gridsearch.best_params_)
```

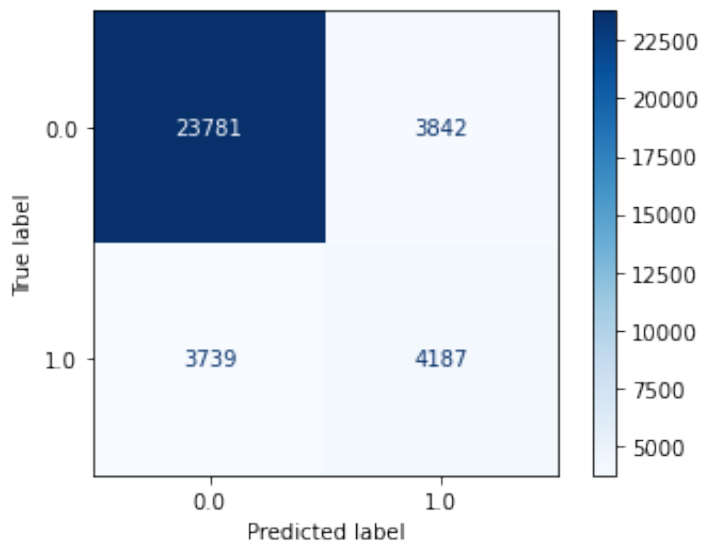
```
{'dt__criterion': 'entropy', 'dt__max_features': None, 'dt__max_leaf_nodes': 5, 'dt__random_state': 42}
```

In [88]:

```
# Create confusion matrix for baseline

cm = plot_confusion_matrix(dt_base, X_test, y_test,
                           cmap=plt.cm.Blues)

plt.grid(False)
plt.savefig('Visualizations/DT.png', bbox_inches = 'tight')
plt.show()
```




In [89]:

```
# Create confusion matrix for best

plot_confusion_matrix(dt_gridsearch, X_test, y_test,
                      cmap=plt.cm.Blues, values_format = '.5g')

plt.grid(False)
plt.savefig('Visualizations/DT.png', bbox_inches = 'tight')
plt.show()
```



As we can see here, the best decision tree model is slightly worse than the logistic regression model, and better than our best kNN model, at a recall score of .747.

## Section 5: Results

### Best Model

The logistic regression model is our best model, at a recall of around .771. Now, I will take a look at the remaining metrics.

In [90]:

```
# Checking our our train and test data metrics

final_test = lr_gridsearch.predict(X_test)
final_train = lr_gridsearch.predict(X_train_resampled)

print("Training Data Results:\n")
print(classification_report(y_train_resampled, final_train))
print("\nTest Data Results:\n")
print(classification_report(y_test, final_test))
```

## Training Data Results:

	precision	recall	f1-score	support
0.0	0.79	0.80	0.79	82693
1.0	0.80	0.78	0.79	82693
accuracy			0.79	165386
macro avg	0.79	0.79	0.79	165386
weighted avg	0.79	0.79	0.79	165386

## Test Data Results:

	precision	recall	f1-score	support
0.0	0.92	0.79	0.85	27623
1.0	0.52	0.77	0.62	7926
accuracy			0.79	35549
macro avg	0.72	0.78	0.74	35549
weighted avg	0.83	0.79	0.80	35549

In [91]:

```
# Creating ROC curve to see effectiveness of model.


y_score = lr_gridsearch.decision_function(X_test)

fpr, tpr, thresholds = roc_curve(y_test, y_score)

sns.set_style('darkgrid', {'axes.facecolor': '0.9'})

print('\nAUC: {}'.format(auc(fpr, tpr)))
plt.figure(figsize=(10, 8))
lw = 2
plt.plot(fpr, tpr, color='darkorange',
         lw=lw, label='ROC curve')
plt.plot([0, 1], [0, 1], color='navy', lw=lw, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.yticks([i/20.0 for i in range(21)])
plt.xticks([i/20.0 for i in range(21)])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.savefig('Visualizations/AUC.png', bbox_inches = 'tight')
plt.show()
```

AUC: 0.8653919990407596




In [92]:

```
# Create confusion matrix for final model
lrcm = confusion_matrix(y_test, final_test)
sns.set_context("talk")
sns.set_theme(style='darkgrid')
ax = sns.heatmap(lrcm, annot=True, cmap='Blues', fmt = 'g', linewidth=0.3, cb

ax.set_title('Results for Next-Day Rain Predictions\n');
ax.set_xlabel('\nPredicted Values')
ax.set_ylabel('Actual Values\n');

## Ticket labels - List must be in alphabetical order
ax.xaxis.set_ticklabels(["Didn't Rain", "Rained"])
ax.yaxis.set_ticklabels(["Didn't Rain", "Rained"])
plt.savefig('Visualizations/FinalCM.png', bbox_inches = 'tight')
## Display the visualization of the Confusion Matrix.
plt.show()
```



**Accuracy:** This model has 79% accuracy, meaning that it correctly determines that it will rain the next day 79% of the time. As True Negatives account for 77% of our test data, there is a lot of bias within this metric. Therefore, it is best to ignore it in our analysis.

**Precision:** If our model says that it will rain tomorrow, there between a 52% chance that it is a true positive and will actually rain the next day. If our model says it wont rain tomorrow, there is a 92% chance that it won't rain tomorrow. The weighted average precision of our model is 83%. Our model is way better at predicting when it won't rain than when it will. In the future, we should take steps to try to raise precision.

**Recall:** Our recall score is the most important aspect of this model. For instances when it actually rained the next day, our model correctly classified that it would rain 77% of the time. For instances when it did not rain the next day, our model correctly predicted that it would not rain 79% of the time. The weighted average recall of our model is 79%.

Our AUC Score is .86. That means that our classifiers have an 86% true positive rate.